



HW 2 (5 late days)

1.



Q1: Training and test accuracy of vanilla ID3




Q1: Number of nodes from vanilla ID3

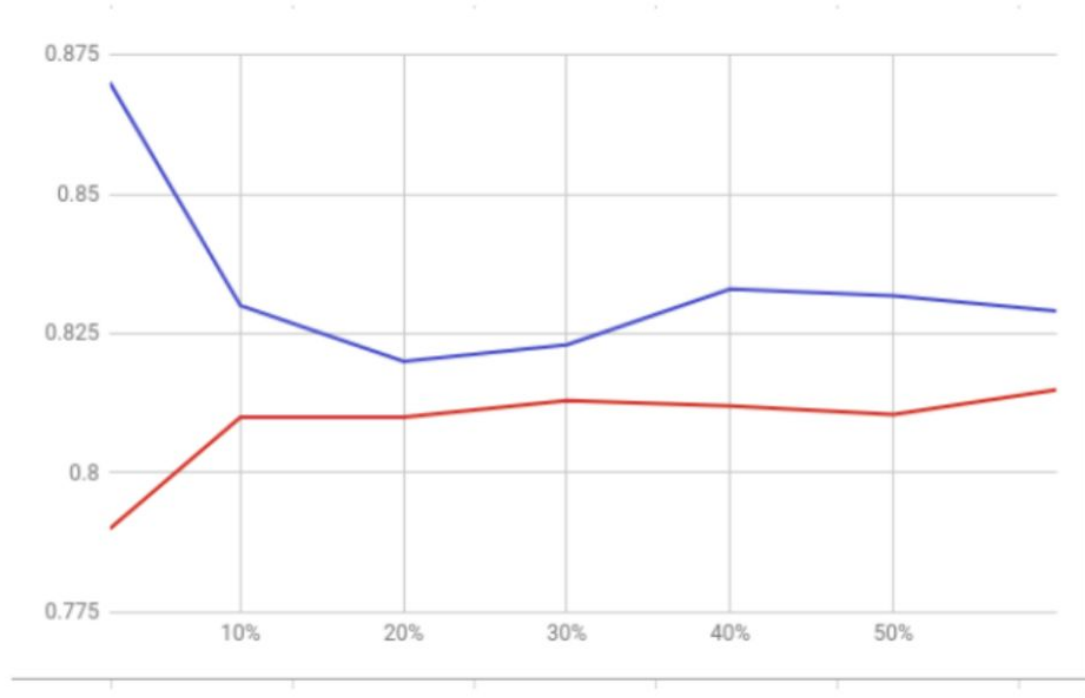



2.





3.





4. The testing set is only supposed to represent unseen data. If we prune it, then we would fit our model to it inadvertently, eventually getting an inaccurate measure of its performance on unseen data.
5. I would keep track of all possible class labels, instead of a single class label, and then output the class labels in order of decreasing frequency.